
Acquiring Grounded Representations of Words with Situated Interactive Instruction

Shiwali Mohan
Aaron H. Mininger
James R. Kirk
John E. Laird

SHIWALI@UMICH.EDU
MININGER@UMICH.EDU
JRKIRK@UMICH.EDU
LAIRD@UMICH.EDU

Computer Science and Engineering, Department of EECS, University of Michigan, Ann Arbor, MI

Abstract

We present an approach for acquiring grounded representations of words from mixed-initiative, situated interactions with a human instructor. The work focuses on the acquisition of diverse types of knowledge including perceptual, semantic, and procedural knowledge along with learning grounded meanings. Interactive learning allows the agent to control its learning by requesting instructions about unknown concepts, making learning efficient. Our approach has been instantiated in Soar and has been evaluated on a table-top robotic arm capable of manipulating small objects.

1. Introduction

A goal of our research is to develop autonomous agents that can dynamically extend not only their knowledge of the world and their ability to interact with it, but also their use of language for interacting with humans. We describe initial progress in developing such an agent that learns grounded representations of adjectives, nouns, prepositions, and verbs through *interactive situated instruction* with a human mentor while performing a task. The agent is developed in Soar (Laird, 2012) using the same architectural mechanisms used in other cognitive tasks without modification.

Our first claim is that in this approach, learning can exploit a broad range of information: knowledge extracted from perception; learned semantic, procedural, and episodic knowledge; and action-model knowledge. Our second claim is that the learning is incremental (learned words aid in learning additional words), online (learning occurs during performance without offline processing), and fast (very few examples are required to teach new words). The agent is reactive; the responses to the human instructor are generated in real time (< 2 seconds). These claims are established by demonstration; we create an agent using this approach that learns via interaction with a human, determine what sources of knowledge it uses in learning, and evaluate the characteristics of its learning.

Our remaining claims focus on the properties of situated interactive instruction. One claim is that situated instructions is an effective method for learning grounded representations of words that combine the words with perceived regularities in the environment. The agent has a vision system for perception and an arm for manipulating objects in its world. For nouns and adjectives, the agent learns new classifications of perceptual features (color, size, and shape) from interactive training with a human and also learns to associate these classifications with specific words (such as *red*, *large*, or *cylinder*). For prepositions, the agent learns to associate combinations of primitive spatial

predicates (such as the alignment of objects) with new words (such as *right of*), and for verbs, the agent learns to associate sequences of primitive actions (such as *pick up* and *put down* with new words (such as *move*). The agent begins with no prior knowledge of these nouns, adjectives, prepositions, and verbs, except for some limited part-of-speech knowledge used in parsing. As the agent learns, the concepts are grounded in its experiences and are dependent on the specifics of the training it receives. For example, the word *red* could be associated with shape instead of color if that is what the instructor wishes. This claim is evaluated by teaching and testing the agent on a variety of nouns/adjectives, prepositions, and verbs.

Our final claim is that interactive instruction is flexible and efficient. By using interactive instructions, an instructor can be freed from using a specific ordering to teach the agent new words and concepts. Often in human controlled interactive learning such as learning by demonstration, the onus is on the instructor to provide good examples from the feature spaces so that the agent can acquire general hypotheses. The instructor must attempt to build and maintain an internal model of what the agent knows and doesn't know. This is especially challenging when the agent is dynamically creating categories from real-world data. In contrast, with mixed initiative interaction, the instructor can rely on the agent to initiate an interaction when needed. This approach can speed instruction by eliminating the need for the instructor to carefully structure the interaction or repeatedly check with the agent to ensure it has completely learned a concept. The agent can actively seek examples of concepts that are hard to learn and avoid asking for multiple examples of easily acquired concepts. The instructor can take initiative in presenting interesting examples to the agent that it might have overlooked, refining agent's learning. We evaluate this claim by demonstrating that the agent initiates appropriate interactions to acquire missing knowledge and comparing the number of examples collected by the agent for different concepts while learning nouns, adjectives, prepositions, and verbs.

The rest of the paper is organized as follows. In section 2, we give an overview of our system including the perceptual and actuation components of our agent. Section 3 provides a brief background of the Soar cognitive architecture emphasizing the mechanisms relevant to our agent design. An overview of the agent design and different phases in learning with instruction are in Section 4. Section 5 describes our human-agent interaction component that forms the basis of learning with instruction. In sections 6, 7, and 8, we describe acquisition of perceptual nouns and adjectives, spatial prepositions, and action verbs. We present the empirical evaluation of the system in 9. We conclude with a description of prior related work in section 10, and the future directions in section 11.

2. System Overview

Our agent exists in a simple table-top environment¹ with a robot arm, Kinect sensor, and four pre-defined locations - a *stove*, a *dishwasher*, a *garbage* and a *pantry*. There are a variety of simple foam blocks of different colors, sizes, and shapes that the arm can manipulate. Figure 1 shows the environment and how the agent interfaces with it.

- **Perception:** The perception system segments the scene into objects using colored 3D point cloud data provided by an overhead Kinect camera. For each of the three perceptual properties, features are extracted and classified using a K-Nearest Neighbor (KNN) classifier with Gaussian weightings. The classification results in a perceptual symbol previously learned by the agent. As

¹ developed by Edwin Olson and his research group at University of Michigan.




Figure 1. System Overview including Major Soar Components

an example, a perceptual symbol R43 is associated with the region in the color feature space that corresponds to the word *red*. These symbols, along with position and bounding box information, are used to create a symbolic representation of the object, which is provided to the agent. The classifiers are trained exclusively through instruction; the process is described in Section 6.

- **Actuation:** To act in the world, the agent sends discrete commands to the robot controller. The current commands include: *point-to(object)*, *pick-up (object)*, *put-down (x,y)*. The robot controller then calculates and performs the motor actions to execute these commands.
- **Instructor Interface:** The instructor interacts with the agent through a simple chat interface. Messages to the agent are fed through the Link-Grammar parser (Sleator & Temperley, 1995) to extract part of speech tags and sentence structure. Messages from the agent are converted to natural language using templates. The instructor can select an object by clicking on it in a live camera feed, and the selection is made known to the agent.

3. Soar Cognitive Architecture

Our agent is implemented in Soar, which has been applied to a wide variety of domains and tasks, including natural language understanding and robot control. Recent extensions to Soar, including episodic and semantic memories, as well as a visual-spatial system, enhance Soar’s ability to support grounded language learning. Relevant components are described in the following paragraphs.

Soar contains a task-independent **Spatial Visual System** (SVS) that supports translations between the continuous representations required for perception and the symbolic, relational representations in Soar. The continuous environment state is represented in SVS as a scene graph composed of discrete objects and their continuous properties. Binary spatial predicates are computed when an agent issues a query for a specific predicate such as *X-axis-aligned(A,B)*. The set of predicates is task independent and fixed, but predicate extraction is controlled using task-specific knowledge.

Working memory maintains symbolic relational representations of current and recent sensory data, current goals, and the agent’s interpretation of the current situation including mappings between objects in the scene and internal symbols and words. Working memory buffers provide interfaces to Soar’s long-term memories, the perception and action systems, and the instructor interface.

Procedural memory contains Soar’s knowledge of how to select and perform actions (called operators), encoded as if-then rules. The locus of decision making is not the selection of a rule. Instead, Soar fires all rules in parallel. The rules propose, evaluate, or apply operators, which are the locus of decision making. Only a single operator can be selected at a time, and once an operator is selected, rules sensitive to its selection and the current context perform its actions (both internal and external) by modifying working memory. Whenever procedural knowledge for selecting or applying an operator is incomplete or in conflict, an impasse occurs and a substate is created in which more

reasoning can occur, including task decomposition, planning, and search methods. In Soar, complex behavior arises not from complex, preprogrammed plans or sequential procedural knowledge, but from the interplay of the agent’s knowledge (or lack thereof) and the dynamics of the environment. In our agent, procedural memory holds rules that implement the processing capabilities such as lexical processing, human-agent interaction, grounded comprehension, and acquisition of grounded representations of words. The agent also has rules that implement the primitive actions and their models. The acquired action-execution knowledge for verbs is stored in procedural memory.

Chunking is a learning mechanism that creates rules from the reasoning that occurred in a substate. When a result is created in a substate, a rule is compiled. The conditions of this rule are the working-memory elements that existed before the substate and were necessary for creating the result, and the actions are the result. The rule is added to procedural memory and is immediately available. Chunking is the mechanism that learns the action-execution knowledge for novel verbs.

Semantic memory stores context independent declarative facts about the world. The agent can store working memory elements into semantic memory and it can retrieve them by creating a cue in a working memory buffer. The best match to the cue (biased by recency and frequency) is retrieved from semantic memory to working memory. In our agent, semantic memory stores *linguistic mapping* knowledge, such as the mapping between a word and a perceptual symbol (red color corresponds to symbol r43). Apart from linguistic mapping knowledge, semantic memory also stores *compositions* of spatial primitives and *action-concept networks* (discussed later). One advantage of semantic memory over procedural memory is that any aspect of a memory can be used for retrieval, whereas in procedural memory, there is an asymmetry between the conditions and actions. An agent can use *red* as a cue, or it could use r43 as a cue, depending on what knowledge is available and what knowledge it needed to retrieve.

Episodic memory stores context-dependent records of the agent’s experiences. It takes snapshots of working memory (episodes) and stores them in chronological order, enabling the agent to retrieve both the context and temporal relations of past experiences. The agent can deliberately retrieve an episode by creating a cue in a working memory buffer. The best partial match (biased by recency) is retrieved and added to working memory. Episodic memory facilitates acquisition of action-execution knowledge through retrospective forward projection by automatically encoding all interactions accompanied by changes in sensory perception. The agent can review past instructions and observe the resulting changes in the environment and its own internal state.

4. Agent Overview

The agent has procedural knowledge for many types of processing including perceptual processing, initiating actions, dialog management, grounded linguistic processing, and learning. During an *interaction cycle*, the agent selects and applies operators to interpret instructor’s utterances and generate behavior. The interaction cycle begins with a natural language utterance from the instructor (shown using solid line in Figure 2) and is processed in following phases.

1. **Lexical Processing:** LG-Soar (Lonsdale et al., 2006), a natural language component implemented as operators in Soar, generates a syntactic parse of the utterance using a static dictionary and grammar. It uses part-of-speech tags to create a parse in the agent’s working memory, identifying the useful content in the message. This parse is further categorized as verb-command, goal-description, descriptive-sentence, etc. based on its lexical structure.

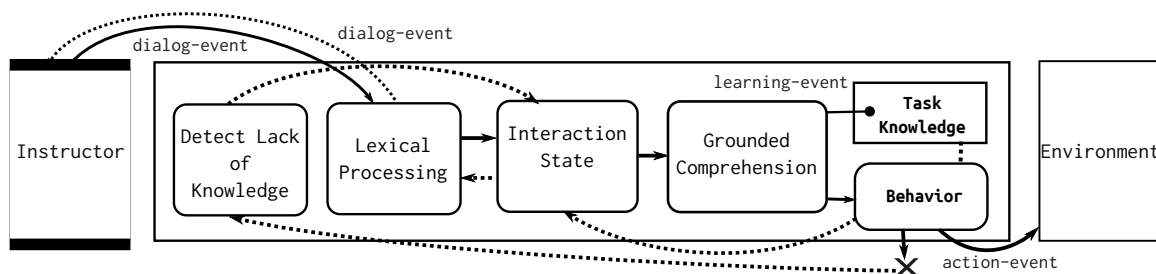


Figure 2. Phases in Interaction Cycle

2. **Interaction Management:** After the utterance has been categorized, it is interpreted within the context of the ongoing dialog between the instructor and the agent. Using the context and the heuristically determined intentions of the utterance, the agent creates a goal to pursue. The goal may include performing actions in the environment in response to commands from the instructor or providing responses to instructor’s queries. Interaction management is done using the mixed-initiative interaction system described in Section 5.
3. **Grounded Comprehension:** To gain useful information from an instruction, the agent must ground linguistic references to objects, spatial relationships, and actions. We use the term *map* for structures in semantic memory that encode how linguistic symbols (nouns/adjectives, spatial prepositions, and action verbs) are associated with perceptual symbols, spatial compositions, and action-concept networks. Maps are learned through interaction with the environment and the instructor and stored in semantic memory. To ground a sentence, the *indexing* process (Mohan & Laird, 2012) attempts to retrieve relevant maps from semantic memory so that it can connect the linguistic terms with their referents. If the terms are successfully mapped, the agent uses constraints derived from the retrieved maps, the current environment, action models, and the interaction context to create a grounded representation of the instruction. These sources of knowledge resolve several ambiguities that can arise from incomplete referring expressions and polysemous verbs in human-agent interactions (Mohan & Laird, 2012). If indexing fails to retrieve a map or there is insufficient knowledge to resolve the ambiguity, an impasse will arise (See Phase 5).
4. **Behavior:** If the agent is successful in generating a grounded representation of the instructor’s utterance, it attempts to pursue the goal associated with the utterance in Phase 2. A natural language command *Pick up the red triangle* results in the agent picking up the referenced object. Apart from actions, the agent can also generate descriptions of the scene and can be queried about various objects and spatial relationships to verify its learning. If there is no action-execution knowledge associated with a command, an impasse will arise (See Phase 5).
5. **Impasse and Acquisition:** If the agent fails to generate a grounded representation or is unable to execute a command, the agent has insufficient knowledge and an impasse arises. In response to the impasse, the agent initiates a new interaction with the instructor (shown using dotted lines in Figure 2) whose purpose is to acquire the missing knowledge. If there are multiple failures during interpretation of a new instruction, the agent processes them one at a time, leading the instructor through a series of interactions until the agent can resolve all the impasses. From an impasse arising in Phase 3 and the ensuing interactions, the agent can learn maps for nouns/adjectives (Section 6), spatial prepositions (Section 7), and action verbs (Section 8) leveraging the structure of interactions. Further semantic information about words including perceptual symbol cate-

gories (Section 6) and spatial compositions (Section 7) are also acquired from impasses arising in Phase 3. Action-concept network and execution knowledge (Section 8) is acquired through interactions initiated due to impasses in the behavior execution phase.

5. Mixed-Initiative Interaction

The agent must maintain a state and context of interactions to comprehend instructions and learn from them, as it is acting in the environment. Thus, the agent needs a model of task-oriented interactive instruction. The interaction model must support the requirements in Table 1.

Table 1. Requirements of Mixed Initiative Interaction Model

| Requirement | Type | Description |
|---------------|-----------------------|---|
| Integrative | I1 | Integrate capabilities, allows agent to plan over and reason about a combined space of linguistic processing, behavior, and learning. |
| Flexible | F1: Initiation | Both instructor and the agent can direct the interaction. |
| | F2: Knowledge | Accommodates communication regarding nouns/adjectives, prepositions, verbs, and related questions. |
| Task-oriented | T1: Contextual | Captures discourse context useful for comprehending incomplete sentences, resolving referent ambiguities, and learning verbs. |
| | T2: Relevant | Instruction-oriented interpretation of human utterances. Agent can ask task-relevant queries. |
| | T3: Structural | Organizes dialog so that it is useful in task execution and learning. |

Our interaction model is based on a theory of discourse structure that stresses the role of purpose in discourse and has been adapted from Rich & Sidner (1998), extending their framework to accommodate learning from situated instruction. The state of interaction is represented by *events*, *segments*, and the *interaction stack*. Figure 3 is an annotated trace of how these concepts are used by the agent while learning the verb *store*. The agent also needs to learn the adjective *orange* and preposition *in* in order to learn the verb.

- **Event:** An event causes change, either in the environment (*action-event*), the discourse state of the instructor-agent interaction (*dialog-event*), or the agent’s knowledge (*learning-event*). Action-events correspond to actions that occur in the environment. Utterances are categorized as dialog-events which are assigned to different classes based on their lexical and syntactic structures, such as *get-next-task*, *verb-command*, and *attribute-query* as shown in the left most column of Figure 3. A learning-event is the successful acquisition of linguistic mapping, semantic, or procedural knowledge. The state space of interactions is defined over utterances, actions, and learning, addressing the *Integrative* (I1) requirement.
- **Segment:** A discourse segment is a contiguous sequence of events that serves a specific *purpose* and organizes a dialog into purpose-oriented interactions. When initiated by the agent, the segments allow the agent to learn a new verb (A1), acquire a mapping for a novel word (such as O11 and P121), acquire a goal (G12), or acquire an action sequence that is required to execute a verb (A13, A14, A15). The segments provide the context for the agent to organize its processing and interactions in pursuit of its task goals (requirement T3). We have encoded heuristics specific to the problem of learning with instruction (requirement T2) that are useful in determination of the purpose (a set of events) of a segment. These heuristics also influence how the agent interprets instructor utterances. Sentences such as *The orange object is in the garbage* are treated as learning

| | | | | |
|----------------|----------------|---|----------|-------|
| get-next-task | dialog-event | Agent: Waiting for the next command... | | |
| verb-command | dialog-event | Human: Store the orange object. | | |
| attribute-q | dialog-event | Agent: What kind of attribute is orange? | | |
| attribute-ans | dialog-event | Human: A color | O11 | |
| | learning-event | Agent: (learns semantic categorization) | | |
| | dialog-event | Agent: I don't see an orange object. Please teach me to recognize one | | |
| | dialog-event | Human: (clicks on the object). This is orange. | O12 | |
| | learning-event | Agent: (learns perceptual symbol + class) | | |
| get-goal | dialog-event | Agent: What is the goal of the action? | | |
| goal-semantics | dialog-event | Human: The goal is the orange object in the pantry. | | |
| relation-q | dialog-event | Agent: I don't know the preposition in. | G12 | |
| | dialog-event | Human: The orange object is in the garbage. | | |
| relation-ans | learning-event | Agent: (learns semantic composition of 'in') | P121 | |
| | learning-event | Agent: (learns goal of 'store') | | |
| | dialog-event | Agent: What action should I take next? | | |
| | dialog-event | Human: Pick up the orange object. | A13 | |
| get-action | action-event | Agent: (picks up the orange triangle) | | |
| verb-command | dialog-event | Agent: What action should I take next? | | |
| | dialog-event | Human: Put the object in the pantry. | A14 | |
| | action-event | Agent: (puts the triangle in the pantry) | | |
| | dialog-event | Agent: What action should I take next? | | |
| | dialog-event | Human: You are done. | A15 | |
| | learning-event | Agent: (learns 'store' with retrospective projection) | | |
| Dialog-class | Events | Dialog | Segments | Stack |

P121: learn preposition in
 G12: acquire goal for store
 A1: external action store

Figure 3. Annotated Human-Agent Dialog for Acquisition of *store*

examples. The segments also encode the reason why they were initiated (requirement T1) which informs language parsing, comprehension, and learning. The context of the segment 011 is useful when parsing the noun-phrase fragment and learning the word *orange*.

- **Stack:** The attentional structure of discourse is captured in a stack of active segments. When a new segment is created, it is pushed onto the stack. The top segment is the focus of the current interaction, and the agent acts to achieve its purpose. When the purpose of the top segment is achieved, it is popped from the stack. The right-most column in Figure 3 shows a snapshot of the stack. It contains three open segments, (P121, G12, and A1), with P121 being the top segment. The segments are hierarchically ordered with each segment contributing towards achieving its parent's purpose (which is lower in the stack). To learn *store*, the agent must acquire a description of the goal and must learn the spatial concept corresponding to the preposition *in*. The stack captures the current context of the dialog (aiding in requirement T1). The model allows any participant to initiate a segment on the stack, supporting the *flexibility* requirement, F1.

6. Perceptual Noun/Adjective Acquisition

A key step in the Grounded Comprehension phase (Section 4) is the resolution of references made to objects in the environment. Our agent can resolve *gestural* references (selecting an object in the camera feed), *descriptive* references (noun phrases like *the large triangle*), and *spatial* references (*the triangle in the pantry*). To ground nouns/adjectives used in descriptive references the agent must connect each word to a set of features in the perceptual system.

For each perceptual property (e.g. color), different features are extracted by the perceptual system to create a separate feature space. For example, the feature space for color is the average RGB values of the object's point cloud. Through instruction, a classifier is trained in each feature space

and it partitions the space into regions, each corresponding to a perceptual symbol used by the agent. The agent must learn both a *linguistic mapping* (the word to a perceptual symbol) and a *perceptual mapping* (the perceptual symbol to a region in a feature space).

6.1 Characterization

Our initial focus is on nouns/adjectives, such as *orange*, *large*, and *triangle*, that describe the perceptual properties: color, size, and shape. Objects are assumed to be monochromatic and so can be described by single colors. Sizes are considered uniform across objects, and relative sizes between objects are not considered. Shapes are limited to those with a distinctive outline from above.

The learning of new nouns and adjectives is *fast*; it requires only a few interactions. This results from several factors, including the characteristics of the objects used, the classification algorithm, the interaction model, and the use of engineered feature spaces for each perceptual property. The use of highly specific feature extractors means we cannot teach perceptual nouns/adjectives not easily captured by one of those feature spaces.

Our approach is also *incremental*. For example, as soon as a color is taught, it can be used to distinguish objects when teaching new verbs or prepositions. Splitting the feature space into specific properties also allows the agent to compose adjectives to recognize objects it has not seen before.

6.2 Background Knowledge

Linguistic: The agent begins with knowledge that *color*, *shape*, and *size* are perceptual properties. The agent does not start any knowledge of specific nouns/adjectives or perceptual symbols.

Perceptual: The perceptual system has pre-encoded knowledge about how to extract useful features for each of the three perceptual properties. Within each property’s feature space, the perceptual system starts with no initial partitioning and no pre-trained data.

6.3 Acquisition

Linguistic: When presented with a new noun/adjective like *orange*, the agent initiates an interaction to learn which of the three perceptual properties it describes. This knowledge is acquired through an explicit interaction with the instructor (Figure 3, 011). The response from the instructor along with the word is used to create a new perceptual symbol (eg. create the symbol *c31* for the color orange). A mapping from the word to the perceptual symbol is stored in semantic memory. We use perceptual symbols instead of the words directly when communicating with the perceptual system in order to add a layer of indirection. This layer gives future flexibility if two words have the same perceptual grounding (synonyms) or one word has multiple groundings (orange as a shape or color).

Perceptual: When the instructor uses a noun or adjective to describe an object (e.g. *This is orange*, Figure 3, 012), the agent uses the linguistic mapping to get a perceptual symbol. It then tells the visual system to use that object as a training example for the corresponding perceptual symbol and its property (e.g. use obj7 as an example of *c14* in the color classifier). This example refines the classifier and the region in the feature space corresponding to the perceptual symbol, thus improving the agent’s perceptual knowledge. Our approach does not depend on any particular type of classifier. A KNN is used because it easily incorporates new training examples, it handles multiple classes, and it allows a perceptual symbol to connect to disjoint and non-linear regions of the feature space.

7. Spatial Preposition Acquisition

In order to naturally interact with objects in the world, the agent needs to understand their relative positions with respect to other objects and locations. To interact with a human mentor, the agent must understand how prepositions, such as *next to*, *in*, or *to the right of*, connect to these spatial relations. Understanding spatial relations and their connection to language is also important to executing commands in the world, where the goal is often to change spatial relations between objects.

The spatial preposition acquisition process supports many modes of learning and interaction with prepositions. The human user can teach a new preposition with a grounded example in the world (*The triangle is left of the square.*), query the system about current relations (*What is left of the square?*), resolve ambiguity (*[Which block?]-The one in the pantry*), use learned prepositions in actions (*Put the square in the pantry.*), and track spatial goals of actions (*The goal is the square in the pantry.*) These capabilities require grounding the linguistic prepositional term to a representation of the appropriate spatial relation as well the ability to use that information for comprehension, communication, and action.

The learning of spatial relationships and their associated prepositions depends on spatial, semantic, and linguistic knowledge. Spatial knowledge, in the form of primitive spatial relationships, is built into the system, while the logical combination of spatial primitives (semantic knowledge) and the mapping of the prepositional term to this combination (linguistic knowledge) are learned. Both the semantic and linguistic knowledge are represented as a network in agent’s semantic memory. In the following sections, we characterize preposition acquisition and give implementation details.

7.1 Characterization

The training example of a preposition requires three components: a primary object, a reference object, and a linguistic term for the preposition. In the teaching statement *the red object is left of the blue object*, the *red object* is the primary object, the *blue object* is the reference object, and *left of* is the preposition. The agent learns only one perspective, the one the mentor is using, and cannot change perspective unless the camera is moved or the prepositions are retaught. The agent can acquire prepositions that are characterized along the following dimensions.

- **Primitives:** The agent can learn prepositions that are compositions of the following primitives.
 - **Direction:** The directional primitives describe how the reference and primary objects are aligned along each axis in a 3-dimensional coordinate system: X, Y, and Z. In relation to the reference object, the primary object can be aligned, greater-than, or less-than. For example two objects that are Z-aligned are on the same plane. These relations are useful in learning prepositions that are based on spatial order, such as *right of* or *diagonal with*.
 - **Distance:** These primitives encode the distance between the reference and the primary object along each axis. The distance is measured from the closest surface of each object. Distance-based primitives are useful in the acquisition of prepositions such as *near* or *far*.
- **Composition:** The learned spatial relations for prepositions are represented by a logical combination of directional primitives and a distribution of distance-based primitives. The combination of directional primitives contain *conjunctions* from different axes, such as X-less-than and Z-aligned, and *disjunctions* on the same axis, such as Y-aligned or Y-greater-than. The initial teaching demonstration results in a representation with a conjunction of the current true directional primitives. Subsequent demonstrations can add disjunctive primitives. Additional




Figure 4. A top down view of four objects in the world, and the associated learned representation of *left of* learned from the three marked examples.

demonstrations provide a distribution of distances from which a range can be calculated. Logical combinations of primitives allow the agent to acquire a wide range of complex spatial prepositions, including ones based on both distance and direction such as *next-to*.

The acquisition process is both fast and general, learning a large variety of useful prepositions from a small number of examples. This learning is also incremental, both in how it uses the knowledge gained from the other learning mechanism, like nouns and adjectives, and in how the representation of spatial primitives is refined with additional training examples.

7.2 Background Knowledge

The directional and distance based primitives are extracted by domain-independent mechanisms encoded in the Spatial Visual System. The agent can query SVS for the value of the primitives for any object pair in the perceptual field.

7.3 Acquisition

The acquisition of spatial prepositions is described below as it relates to the two types of learned knowledge: semantic and linguistic. The mentor teaches the agent a new spatial relation by referring to two objects in the world that demonstrate that relationship. The teaching example *The square is left of the circle* is explored in the following sections and depicted in Figure 4.

Semantic: After the sentence is parsed and the objects indexed, SVS is queried for all of the current directional primitive relations between the specified objects and the values of the distance primitives. In Figure 4, three successive teaching examples are shown; the reference object is the blue circle, and the three numbered squares are the primary objects for the three examples. The first teaching example exhibits the directional primitives X-less-than, Z-aligned, and Y-aligned. This enables the acquisition of a specific sense of *left of*, with a strict alignment, with just one example. However it is possible that the learned information is more specific than intended. The second example teaches that Y-aligned or Y-less-than is permissible and the third example teaches that Y-greater-than is acceptable as well, forming a disjunction of primitives for the Y axis. With these additional examples, the agent has a distribution of distances along the X and Y axes that are used to calculate a range. This compositional structure, including the distance ranges and conjunctions and disjunctions of directional primitives, is stored in semantic memory. The right-side of the figure displays a representation of this final version of this structure.

Distances from one example are insufficient to learn whether a distance-based metric is related to the preposition. Given multiple examples, the distribution of distances is used to determine an acceptable distance range for the relation. For a relation like *near* this range may be very restrictive, but for *left of*, a large range indicates that distance is not an important metric for that preposition.

The system also records if any examples are aligned on every axis simultaneously so that the system can distinguish examples that are explicitly inside or intersecting.

Linguistic Mapping: The preposition linguistic term used to describe the new relationship is mapped to the learned compositional structure in semantic memory. In the above example, *left of* is mapped to this representation and stored into semantic memory, where it can be accessed when *left of* is used in the future.

7.4 Spatial Projection

So far we have described how the system acquires and represents spatial prepositions. This knowledge can be accessed with a lookup into semantic memory using the linguistic prepositional term whenever the system needs to know how two objects spatially relate. When executing an action in the world that attempts to establish the spatial relation underlying a preposition (such as *put the object to the left of the pantry*), the system needs to be able to *project* that learned preposition to a specific point (X, Y, Z) in the world.

The current representation provides sufficient information to project to a point in the world. The legal alignments along each axis are known as well as the average distance of the examples. When there are multiple legal alignments for one axis, one is randomly chosen. If the chosen option is aligned, then a distance of 0 is used and otherwise the average distance. For example take the situation when projecting *put the object to the left of the pantry*, using the learned representation from Figure 4 of *left of*. If the pantry is located at (x_i, y_j, z_k) , the projection point would be $(x_i - 1.7, y_j, z_k)$, representing alignment on the Y and Z axes and the average displacement on the X axis.

8. Action Verb Acquisition

The verb acquisition problem can be decomposed as the acquisition of knowledge in three distinct but related categories. *Linguistic mapping* knowledge allows the agent to associate the linguistic forms (verbs, prepositions, and noun phrases) in an action command with an action-concept network. The agent also acquires these action-concept networks (*semantic knowledge*), which represent a novel action, its parametrization and goals, and their mutual constraints. The linguistic and semantic knowledge is represented in semantic memory. Finally, the agent acquires *procedural knowledge* of the novel action, a composition of the primitive actions that leads to the goal. This knowledge is represented as rules that encode the action selections necessary to achieve the goal of the verb. After successful acquisition, the agent can comprehend and execute verb commands by grounding linguistic forms to objects and actions.

A typical example of the types of verbs that our system can learn is *move*, which is a composition of known primitive action verbs *pick up* and *put down*. On learning the verb *move*, the agent can comprehend the action command *move the red triangle to the pantry* by grounding the arguments *the red triangle* and *the pantry* to objects present in the environment and instantiating (and later, executing) the action that results in *the red triangle in the pantry*. Since, the emphasis of this work is on acquisition of knowledge through interactive instruction, we rely on the human instructor to explicitly provide examples of all knowledge categories described above through natural language.

8.1 Characterization

The agent can acquire action verbs that can be characterized along the following dimensions.

- **Parametrization:** A single verb word may map to different action-concepts based on its argument structure. Some verb commands such as *move the red block to the pantry* make the argument structure *explicit* by including the object and the location in the command. However, in some verb commands such as *store the red block*, an argument (location = pantry) might be *implicit* in the verb *store* itself. The agent can learn both categories of verbs.
- **Goal:** The learning is currently limited to verbs that have perceptible, spatial goal states, such as *the red block is to the right of the pantry*.
- **Composition:** The agent can learn verbs that are compositions of known primitive actions, such as *move* which can be executed by *picking-up* an object and *putting-down* at a specific position. The agent can learn arbitrary long sequences of primitives, if they are useful in achieving the goal.

We are interested in *fast* and *general* learning from a small number of examples of action execution. The agent uses *chunking*, similar to *explanation-based generalization* (Rosenbloom & Laird, 1986) to develop casual connections between primitive actions and the goal of the verb. It extracts rules that are generally applicable from specific instances of action execution.

8.2 Background Knowledge

Linguistic Mapping: To ground an action command the agent must associate the verb and its argument structure to actions and objects in the environment. This mapping for known primitive verbs is encoded declaratively in the agent’s semantic memory and allows the agent to access the related action operator which can be instantiated with objects in the environment. Consider the example shown in Figure 5, Network A. It maps the verb *put* with an argument structure consisting of a *direct-object* and the object connected to the verb via the preposition *in* with the operator *op_put-down*. This allows the agent to associate the sentence *Put a large red block in the table* with an instantiated operator which will achieve the intended goal.

Procedural: The agent has pre-programmed rules that allow it to execute the primitive actions in the environment. The actions are implemented as operators in Soar. An action is defined by its availability conditions (the preconditions of the action), execution knowledge (rules that execute action commands in the environment), and termination conditions (a set of predicates that signify that the goal of the action has been achieved). The agent maintains a set of all available primitive actions given the current physical constraints, object affordances, and the agent’s domain knowledge. The agent also has domain action models (encoded as rules) with which it can simulate the effect of its actions on the environment during learning.

8.3 Acquisition

Acquisition of new verbs is integrated with interactive execution of tasks. If the agent cannot ground an verb command to a known action, it tries to learn it through agent-initiated situated interactions. We use the example in Figure 3 to demonstrate acquisition of the verb *store*.

Linguistic Mapping: Using the argument structure of the action command *Store the orange triangle* (extracted by the syntactic parser), the agent creates a new mapping in its semantic memory (shown in Figure 5, Network B, nodes (M1, L1, A11, P1)). This mapping associates the novel verb *store* and its *direct-object* to a new operator (*op_1*) and its argument. A11 is a slot that can be filled by an object

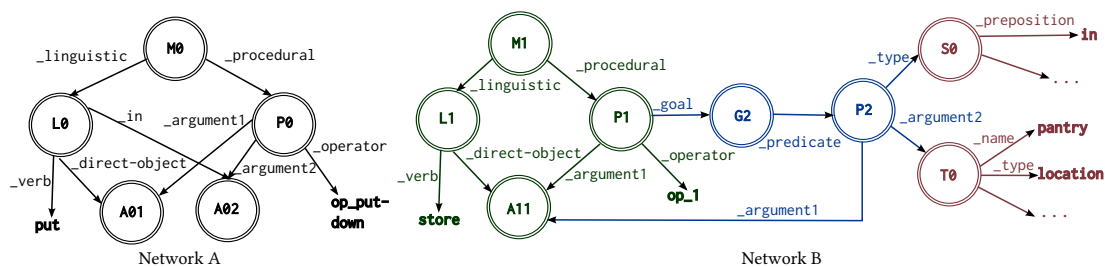


Figure 5. A. Pre-encoded linguistic map for *put*. B. Acquired action-concept network for *store*

that satisfies the description (noun-phrase) connected to the verb as a *direct-object*. corresponds to the *direct-object* of the verb. The edge A11, P1 constrains the instantiation of argument1 to the same object. Future action commands containing the verb *move* can be indexed to the action op_1 using this map. After acquiring this mapping, the agent proposes the operator op_1. However, since it does not know how to execute this operator yet, it experiences an impasse. This impasse causes the agent to begin a new interaction focusing on learning the verb.

Semantic: The agent prompts the instructor for an explicit description of the goal (Figure 3, segment G12). On getting a reply, the agent grounds the utterance to objects and their spatial relationships. This representation is used to augment the action-concept network (Figure 5, nodes (G2, P2)). The edge P2, A11 introduces constraints on how the goal can be defined on the basis of instantiation of the verb. This network creates connections between the linguistic knowledge of the verb and the compositional structure of the preposition *in* and the location *pantry*. These concepts may have been acquired in previous experiences or may be learned with the verb acquisition (as in Figure 3).

Procedural: The agent has not yet acquired execution rules for the verb *store*, hence, the impasse is not resolved. The agent initiates an interaction with the instructor to acquire a situated example execution of the action. Through this interaction, the instructor decomposes the action into a sequence of primitive actions (Figure 3, segments A13, A14), which the agent executes in the environment as they are provided until the goal of the verb is achieved. These interactions are automatically stored in the agent’s episodic memory and are used later for retrospective learning (Mohan et al., 2012) .

Once the goal is achieved, the agent attempts to learn the conditions under which it should execute the instructed primitive action to achieve the goal. This process involves using episodic memory to reconstruct the state that the agent was in when the new verb was first suggested. From there, the agent continues to use episodic memory to direct the selection of the instructed actions, and uses its action-models to simulate the instructed actions. This internal projection generates an explanation from the initial state to the goal. From this explanation, the agent extracts the casual dependencies (via chunking) and learns rules for selecting each one of the component actions.

9. Evaluation

We have described an agent that acquires diverse knowledge including perceptual classification, spatial composition, and procedural and semantic knowledge by grounding natural language instructions to the immediate environment. This establishes our claim that our agent implemented in Soar supports learning from various sources of information including the human instructor. The interaction system provides *flexibility* to support various levels of instructor/agent control over learning. If the agent is unable to progress, it poses queries to the instructor and incorporates the replies in

its domain knowledge ([video](#)¹). The instructor can guide learning by presenting the useful concepts before the agent is asked to perform a command ([video](#)²).

For evaluating our other claims³, we created a space of examples for each category of concepts (noun/adjectives, prepositions, verbs) the agent can acquire. A random ordering (*trial*) of examples was generated. The examples from a trial were presented sequentially to the agent in an *interleaved* training and testing session. The agent was tested for correctness on every example. In case the example was unknown to the agent or the agent was incorrect, the example was added to the training set and relevant instructions were provided. Trials were repeated until the performance converged (two successive trials with 100% performance). We report the results using the average number of examples required to learn the concepts. The results were averaged across 3 runs.

For noun/adjective learning, separate example spaces for properties color, size, and shape were generated. Each space contained 12 objects selected randomly from a set that had 4 distinct colors, 2 sizes, and 4 shapes. The agent was presented with an object and asked for a word associated with the object for that property. For example, the agent was asked *What color is this?* for color. The correct word was provided if the answer was unknown or incorrect. The agent associated correct colors to 100% of the objects on observing an average of 1 example per color. 100% performance was achieved after the agent observed an average of 1.5 examples per size. The agent associated correct shapes to 96% of the objects on observing an average of 12.9 examples per shape.

For evaluating preposition learning, 6 prepositions - *left*, *right*, *in*, *near*, and *far* were selected. For each preposition, two objects (*red*, *blue*) were arranged in a spatial configuration that was thought to be representative for the preposition. The agent was then asked an evaluation question, *Is the blue object to the right of the red object?*. The agent was corrected if it replied incorrectly or reported that the preposition was unknown. The agent was able to correctly recognize 93.98% of 144 spatial arrangement and preposition associations after observing an average of 3.17 examples per preposition.

To generate the space for verb acquisition, we created action command templates from 3 novel verbs, 3 for the verb *move* by combining it with prepositions - *in*, *left of*, and *right of* and 1 each for *store* and *discard*. Templates were instantiated with randomly selected objects (from a set of 4) and locations (a set of 4) to generate commands. The initial state of the arm (holding(obj1) or empty) was also randomly assigned. If the agent asked for instructions, superfluous actions were randomly introduced in the instructions. The training converged after an average of 1.26 examples per action command. We then generated 5 random instantiations of every action template to test the agent. The agent executed all instantiations correctly. To understand the generality in learning, consider the template *move* <obj1> *to the right of* <obj2/location> for which there are 64 possible instantiations; with only 2 examples, the agent can correctly execute all instantiations.

The data presented above shows that the agent’s learning is *fast*. For various concepts, the agent was able to extract general knowledge that covered 93.8% - 100% of the example space from very few (1-13) examples. The interactive learning paradigm allows the agent to be more discriminative while gathering examples, it gathers few examples for concepts that are easy to learn but gathers more examples for concepts that are hard to acquire, making acquisition *efficient*. In noun/adjectives, the agent gathers more examples for shape than for color or size, as shape is harder to learn given the perceptual features the agent uses. For learning prepositions, the agent adapts interactions to gather more examples for concepts that are hard such as *near* for which the agent collects 4.67 examples

¹ http://youtu.be/_ktny-h0KX4 ² <http://youtu.be/9M-rpdXFbgs> ³ conducted by Soar Technology, Ann Arbor




Figure 6. (Left) Number of agent-initiated interactions per action-command. (Right) Agent response time

on an average. This can be compared to the preposition *behind* (reference object is *directly* behind the primary object) for which the agent collects only 1 example. Similarly, in verb acquisition, the agent asks for instructions only when its knowledge is insufficient for making further progress.

Figure 6 (left) shows the agent’s performance while it was learning over a combined space of 9 nouns and adjectives, 3 prepositions, and 3 actions. The example space contained action-commands composed from combining randomly selected verbs, prepositions, and objects. The agent begins with no prior knowledge. In the first command, the agent initiates 24 interactions as it attempts to learn about nouns/adjectives, preposition, and verbs. As more commands are given, the number of interactions reduce as the agent generalizes knowledge gathered from previous commands and applies it in novel situations. Finally, the number of interactions for a command reduce to 1 - the utterance from the instructor. This establishes our claim that agent’s learning is *online* and *incremental*. Figure 6 (right) plots the time taken by the agent to generate responses to instructor’s utterance. The responses that involved internal processing were generated under 1.1 seconds. Responses that involve taking an external action in the environment and learning from projection take longer, but can be interrupted and do not impact the reactivity of the agent.

10. Related Work

Allen et al. (2007) demonstrate a collaborative task learning agent that acquires procedural knowledge through a collaborative session of demonstration, learning, and dialog. The human teacher provides a set of tutorial instructions accompanied with related demonstrations in a shared environment. The agent uses these situated examples to acquire new procedural knowledge. Although the learning is human demonstration driven, the agent controls certain aspects of its learning by making generalizations without requiring the human to provide a large number of examples. A key distinction of our work is that the initiative of learning is placed with both the instructor and the agent. The agent can initiate a learning interaction if its knowledge is insufficient for further progress, and the instructor can verify the agent’s knowledge by asking relevant questions.

Chen & Mooney (2011) describe a unified agent architecture for human-robot collaboration that combines natural language processing and common sense reasoning. The developed a planning agent that relies on communication with the human to acquire further information about under-specified tasks. The agent also demonstrates limited learning by acquiring novel common sense rules through dialog. Our agent uses interactive instruction to learn a wider variety of knowledge, including grounded representations of language including task execution knowledge.

A significant line of research for grounded language acquisition is being pursued by Tellex et al. (2011). Their framework instantiates a probabilistic graphical model for a particular natural

language command according to the command’s hierarchical and compositional structure. They focus on corpus-based learning mechanisms, while we concentrate on interactive learning where the agent learns from a two way communication with the human instructor.

11. Future Work

There are several avenues that we are interested in exploring. The primary focus will be on extending the learning capabilities of the agent. This will involve increasing the complexity of the environment, including more complicated objects, to enable learning more complex and diverse adjectives, verbs, and prepositions.

For nouns and adjectives, we will increase the number and types of perceptual properties we can learn, including non-visual ones like weight and more complicated shapes. The next step is acquiring knowledge for perceptual and functional categorization and classification. We expect the agent to acquire knowledge such as *cans are gray cylinders*, or that only some objects can be picked up. This knowledge could aid in acquiring and generalizing action affordances.

The acquisition of prepositions is currently limited to simple binary spatial relationships between two objects, such as *right of*. We plan to extend the system to acquire more complex relations involving multiple objects, such as *between*, to learn contact based relations, such as *on top*, and to learn logical combinations of objects and prepositions, such as *clear means all objects outside*.

Future extensions will allow the agent to learn the availability conditions for novel actions along with acquisition of execution and termination conditions (discussed previously). Under the current constraints, the agent can acquire action verbs that are compositions of known primitives and have perceptible goals. The goals have to be communicated declaratively to the agent. Our future efforts will focus on acquiring goals autonomously from multiple executions of the verb in the environment. We are interested in expanding the kinds of verbs the agent can learn including non-action verbs.

References

- Allen, J., Chambers, N., Ferguson, G., & Galescu, L. (2007). Plow: A Collaborative Task Learning Agent. In *Proceedings of the National Conference on Artificial Intelligence*. Vancouver, BC.
- Chen, D. L. & Mooney, R. J. (2011). Learning to Interpret Natural Language Navigation Instructions from Observations. In *Proceedings of the Twenty-Fifth AAAI Conference on Artificial Intelligence*, 859–865.
- Laird, J. E. (2012). *The Soar Cognitive Architecture*. MIT Press.
- Lonsdale, D. W., Tustison, C., Parker, C., & Embley, D. (2006). Formulating queries for assessing clinical trial eligibility. In *Proceedings of the 11th International Conference on Applications of Natural Language to Information Systems*.
- Mohan, S. & Laird, J. (2012). Situated Comprehension of Imperative Sentences in Embodied, Cognitive Agents. In *Technical Reports from AAAI Workshop on Grounded Language for Physical Systems*.
- Mohan, S., Mininger, A., Kirk, J., & Laird, J. (2012). Learning Grounded Language through Situated Interactive Instruction. In *Technical Reports from AAAI Fall Symposium Series (to appear)*.
- Rich, C. & Sidner, C. (1998). COLLAGEN: A Collaboration Manager for Software Interface Agents. *User Modeling and User-Adapted Interaction*, 8(3), 315–350.
- Rosenbloom, P. S. & Laird, J. E. (1986). Mapping Explanation-Based Generalizations onto Soar. In *Proceedings of the AAAI conference on Artificial Intelligence*, 561–567.
- Sleator, D. & Temperley, D. (1995). Parsing English with a Link Grammar. *CoRR*.
- Tellex, S., Kollar, T., Dickerson, S., & Walter, M. (2011). Understanding Natural Language Commands for Robotic Navigation and Mobile Manipulation. In *Proceedings of AAAI*.